

# elevvo-ml-intern

August 14, 2025

[https://colab.research.google.com/drive/1WgFGPhsLwmuWpaHR-Pj3X\\_QbWERf1Nff?usp=sharing](https://colab.research.google.com/drive/1WgFGPhsLwmuWpaHR-Pj3X_QbWERf1Nff?usp=sharing)

## Github Repository

```
[6]: %cd https://github.com/zoya4477/ML-intern.git
!git clone
!git config --global user.email "zoyahafeez785@gmail.com"
!git config --global user.name "zoya4477"
```

```
[Errno 2] No such file or directory: 'https://github.com/zoya4477/ML-intern.git'
/content
```

```
fatal: You must specify a repository to clone.
```

```
usage: git clone [<options>] [--] <repo> [<dir>]
```

-v, --verbose	be more verbose
-q, --quiet	be more quiet
--progress	force progress reporting
--reject-shallow	don't clone shallow repository
-n, --no-checkout	don't create a checkout
--bare	create a bare repository
--mirror	create a mirror repository (implies bare)
-l, --local	to clone from a local repository
--no-hardlinks	don't use local hardlinks, always copy
-s, --shared	setup as shared repository
--recurse-submodules[=<pathspec>]	initialize submodules in the clone
--recursive ...	alias of --recurse-submodules
-j, --jobs <n>	number of submodules cloned in parallel
--template <template-directory>	directory from which templates will be used
--reference <repo>	reference repository
--reference-if-able <repo>	reference repository
--dissociate	use --reference only while cloning
-o, --origin <name>	use <name> instead of 'origin' to track upstream
-b, --branch <branch>	checkout <branch> instead of the remote's HEAD
-u, --upload-pack <path>	

```

                                path to git-upload-pack on the remote
--depth <depth>                create a shallow clone of that depth
--shallow-since <time>         create a shallow clone since a specific time
--shallow-exclude <revision>   deepen history of shallow clone, excluding rev
--single-branch                clone only one branch, HEAD or --branch
--no-tags                      don't clone any tags, and make later fetches not to
follow them
--shallow-submodules           any cloned submodules will be shallow
--separate-git-dir <gitdir>    separate git dir from working tree
-c, --config <key=value>      set config inside the new repository
--server-option <server-specific>
                                option to transmit
-4, --ipv4                     use IPv4 addresses only
-6, --ipv6                     use IPv6 addresses only
--filter <args>               object filtering
--remote-submodules           any cloned submodules will use their remote-tracking
branch
--sparse                       initialize sparse-checkout file to include only files
at root

```

## 1 Task 1: Student Score Prediction

```

[ ]: #Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
import seaborn as sns

```

```

[ ]: #Load Dataset
df = pd.read_csv("/content/archive.zip")
df.head()

```

```

[ ]:
Hours_Studied  Attendance  Parental_Involvement  Access_to_Resources  \
0              23          84                    Low             High
1              19          64                    Low             Medium
2              24          98                   Medium             Medium
3              29          89                    Low             Medium
4              19          92                   Medium             Medium

```

	Extracurricular_Activities	Sleep_Hours	Previous_Scores	Motivation_Level \
0	No	7	73	Low
1	No	8	59	Low
2	Yes	7	91	Medium
3	Yes	8	98	Medium
4	Yes	6	65	Medium

	Internet_Access	Tutoring_Sessions	Family_Income	Teacher_Quality \
0	Yes	0	Low	Medium
1	Yes	2	Medium	Medium
2	Yes	2	Medium	Medium
3	Yes	1	Medium	Medium
4	Yes	3	Medium	High

	School_Type	Peer_Influence	Physical_Activity	Learning_Disabilities \
0	Public	Positive	3	No
1	Public	Negative	4	No
2	Public	Neutral	4	No
3	Public	Negative	4	No
4	Public	Neutral	4	No

	Parental_Education_Level	Distance_from_Home	Gender	Exam_Score
0	High School	Near	Male	67
1	College	Moderate	Female	61
2	Postgraduate	Near	Male	74
3	High School	Moderate	Male	71
4	College	Near	Female	70

## Data Cleaning

```
[ ]: # Check for missing values
print(df.isnull().sum())

# Drop rows with nulls or fill them
df = df.dropna() # or df.fillna(method='ffill', inplace=True)

# View data types and check unique values
df.info()
```

Hours_Studied	0
Attendance	0
Parental_Involvement	0
Access_to_Resources	0
Extracurricular_Activities	0
Sleep_Hours	0
Previous_Scores	0
Motivation_Level	0

```

Internet_Access          0
Tutoring_Sessions        0
Family_Income            0
Teacher_Quality          78
School_Type              0
Peer_Influence           0
Physical_Activity        0
Learning_Disabilities     0
Parental_Education_Level 90
Distance_from_Home       67
Gender                   0
Exam_Score               0
dtype: int64
<class 'pandas.core.frame.DataFrame'>
Index: 6378 entries, 0 to 6606
Data columns (total 20 columns):

```

#	Column	Non-Null Count	Dtype
0	Hours_Studied	6378 non-null	int64
1	Attendance	6378 non-null	int64
2	Parental_Involvement	6378 non-null	object
3	Access_to_Resources	6378 non-null	object
4	Extracurricular_Activities	6378 non-null	object
5	Sleep_Hours	6378 non-null	int64
6	Previous_Scores	6378 non-null	int64
7	Motivation_Level	6378 non-null	object
8	Internet_Access	6378 non-null	object
9	Tutoring_Sessions	6378 non-null	int64
10	Family_Income	6378 non-null	object
11	Teacher_Quality	6378 non-null	object
12	School_Type	6378 non-null	object
13	Peer_Influence	6378 non-null	object
14	Physical_Activity	6378 non-null	int64
15	Learning_Disabilities	6378 non-null	object
16	Parental_Education_Level	6378 non-null	object
17	Distance_from_Home	6378 non-null	object
18	Gender	6378 non-null	object
19	Exam_Score	6378 non-null	int64

```

dtypes: int64(7), object(13)
memory usage: 1.0+ MB

```

Basic Data Visualization


```


[ ]: # Correlation matrix
plt.figure(figsize=(10, 8))
# Select only numerical columns for correlation calculation
numerical_df = df.select_dtypes(include=[np.number])
sns.heatmap(numerical_df.corr(), annot=True, cmap='coolwarm')

```

```
plt.title('Correlation Matrix')
plt.show()

# Scatter plot: study time vs final grade
plt.scatter(df['Hours_Studied'], df['Exam_Score'])
plt.xlabel('Hours Studied')
plt.ylabel('Exam Score')
plt.title('Hours Studied vs Exam Score')
plt.show()
```





Feature Selection

```
[ ]: # Let's use Hours_Studied to predict Exam_Score
X = df[['Hours_Studied']] # Features
y = df['Exam_Score']     # Target variable
```

Train Test Split

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```


Train Linear Regression Model

```
[ ]: lr_model = LinearRegression()
lr_model.fit(X_train, y_train)
y_pred = lr_model.predict(X_test)
```

Visualize Prediction

```
[ ]: plt.scatter(X_test, y_test, color='blue', label='Actual')
plt.scatter(X_test, y_pred, color='red', label='Predicted')
```

```
plt.xlabel('Study Time')
plt.ylabel('Final Grade (G3)')
plt.legend()
plt.title('Actual vs Predicted Scores')
plt.show()
```



Model Evaluation

```
[ ]: mse = mean_squared_error(y_test, y_pred)
      r2 = r2_score(y_test, y_pred)

      print(f"Mean Squared Error: {mse:.2f}")
      print(f"R-squared: {r2:.2f}")
```

Mean Squared Error: 12.35

R-squared: 0.21

Polynomial Regression

```
[ ]: poly = PolynomialFeatures(degree=3)
      X_poly = poly.fit_transform(X)
```

```

X_train_p, X_test_p, y_train_p, y_test_p = train_test_split(X_poly, y,
    ↪test_size=0.2, random_state=42)

poly_model = LinearRegression()
poly_model.fit(X_train_p, y_train_p)
y_pred_poly = poly_model.predict(X_test_p)

# Evaluation
mse_poly = mean_squared_error(y_test_p, y_pred_poly)
r2_poly = r2_score(y_test_p, y_pred_poly)

print(f"Polynomial MSE: {mse_poly:.2f}")
print(f"Polynomial R²: {r2_poly:.2f}")

```

Polynomial MSE: 12.37

Polynomial R²: 0.20

Feature Engineering (Multiple Regression)

Try experimenting with different feature combinations (e.g., removing or adding features like sleep, participation, etc.)

```

[ ]: # Experimenting with different feature combinations
features_experiment = ['Hours_Studied', 'Attendance', 'Previous_Scores',
    ↪'Sleep_Hours', 'Physical_Activity']
X_experiment = df[features_experiment]
y_experiment = df['Exam_Score']

# Split the data with the new features
X_train_exp, X_test_exp, y_train_exp, y_test_exp =
    ↪train_test_split(X_experiment, y_experiment, test_size=0.2, random_state=42)

# Train a new linear regression model with the experimental features
lr_model_exp = LinearRegression()
lr_model_exp.fit(X_train_exp, y_train_exp)
y_pred_exp = lr_model_exp.predict(X_test_exp)

# Evaluate the new model
mse_exp = mean_squared_error(y_test_exp, y_pred_exp)
r2_exp = r2_score(y_test_exp, y_pred_exp)

print(f"Experimental Model MSE: {mse_exp:.2f}")
print(f"Experimental Model R-squared: {r2_exp:.2f}")

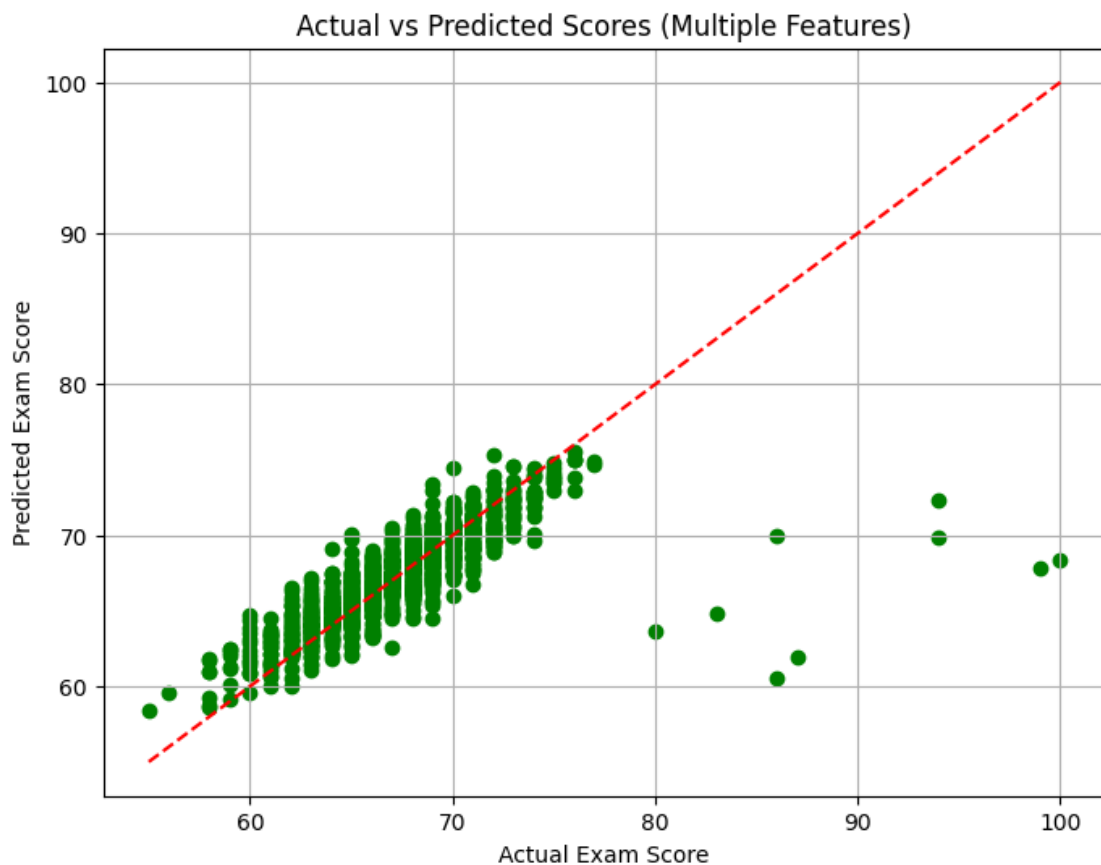
```

Experimental Model MSE: 6.47

Experimental Model R-squared: 0.58



```
[ ]: #Visualize Predictions
plt.figure(figsize=(8, 6))
plt.scatter(y_test_exp, y_pred_exp, color='green')
plt.plot([y_test_exp.min(), y_test_exp.max()], [y_test_exp.min(), y_test_exp.
    ↪max()], '--r')
plt.xlabel("Actual Exam Score")
plt.ylabel("Predicted Exam Score")
plt.title("Actual vs Predicted Scores (Multiple Features)")
plt.grid(True)
plt.show()
```



```
[ ]:
```

## 2 Task 2: Customer Segmentation

```
[ ]: #libraries
import pandas as pd
import numpy as np
```

```
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import silhouette_score
from sklearn.cluster import DBSCAN
```

```
[ ]: #Data Loading
df = pd.read_csv('/content/mall_customer.zip')
df.head()
```

```
[ ]:   CustomerID  Gender  Age  Annual Income (k$)  Spending Score (1-100)
0           1    Male   19             15             39
1           2    Male   21             15             81
2           3  Female   20             16              6
3           4  Female   23             16             77
4           5  Female   31             17             40
```

Data Cleaning

```
[ ]: print(df.info())
print(df.isnull().sum())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CustomerID                            200 non-null   int64
1   Gender                                200 non-null   object
2   Age                                    200 non-null   int64
3   Annual Income (k$)                    200 non-null   int64
4   Spending Score (1-100)                 200 non-null   int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB
None
CustomerID      0
Gender          0
Age             0
Annual Income (k$)  0
Spending Score (1-100)  0
dtype: int64
```


```
[ ]: X = df[['Annual Income (k$)', 'Spending Score (1-100)']]
```

Feature Scaling

```
[ ]: scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

Visualize Raw Data

```
[ ]: plt.figure(figsize=(8, 6))
plt.scatter(X['Annual Income (k$)'], X['Spending Score (1-100)'], c='blue')
plt.xlabel('Annual Income (k$)')
plt.ylabel('Spending Score')
plt.title('Customer Distribution')
plt.show()
```




Optimal K using Elbow Method

```
[ ]: wcss = []
K_range = range(1, 11)

for k in K_range:
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(X_scaled)
```

```
wcss.append(kmeans.inertia_)

plt.plot(K_range, wcss, 'bo-')
plt.xlabel('Number of Clusters (k)')
plt.ylabel('WCSS')
plt.title('Elbow Method For Optimal k')
plt.grid(True)
plt.show()
```




K- Means Clustering

```
[ ]: kmeans = KMeans(n_clusters=5, random_state=42)
df['Cluster'] = kmeans.fit_predict(X_scaled)
```

Visualize Cluster

```
[ ]: plt.figure(figsize=(8, 6))
sns.scatterplot(data=df, x='Annual Income (k$)', y='Spending Score (1-100)',
               hue='Cluster', palette='Set1')
plt.title('Customer Segments (K-Means)')
```

```
plt.legend()
plt.grid(True)
plt.show()
```



Analyze Clusters

```
[ ]: cluster_summary = df.groupby('Cluster')[['Annual Income (k$)', 'Spending Score_
      ↪(1-100)']].mean()
print(" Average values per cluster:\n")
print(cluster_summary)
```


Average values per cluster:

Cluster	Annual Income (k\$)	Spending Score (1-100)
0	55.296296	49.518519
1	86.538462	82.128205
2	25.727273	79.363636
3	88.200000	17.114286
4	26.304348	20.913043

Try DBSCAN

```
[ ]: dbscan = DBSCAN(eps=0.5, min_samples=5)
df['DBSCAN_Cluster'] = dbscan.fit_predict(X_scaled)

plt.figure(figsize=(8, 6))
sns.scatterplot(data=df, x='Annual Income (k$)', y='Spending Score (1-100)',
               hue='DBSCAN_Cluster', palette='Set2')
plt.title('DBSCAN Clustering Results')
plt.grid(True)
plt.show()
```



Analyze Average Spending per DBSCAN Cluster


```
[ ]: # Group by DBSCAN cluster and compute means
dbscan_summary = df.groupby('DBSCAN_Cluster')[['Annual Income (k$)', 'Spending_
Score (1-100)']].mean().round(2)

print(" Average values per DBSCAN cluster:\n")
print(dbscan_summary)
```

Average values per DBSCAN cluster:

DBSCAN_Cluster	Annual Income (k\$)	Spending Score (1-100)
-1	122.75	46.88
0	52.49	43.10
1	82.54	82.80

```
[ ]:
```



### 3 Task 3: Forest Cover Type Classification

```
[ ]:
```

```
#Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import classification_report, confusion_matrix, \
    ConfusionMatrixDisplay
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.preprocessing import StandardScaler
```

```
[ ]:
```

```
#Load Dataset
df = pd.read_csv('/content/Forest_cover.zip')
df.head()
```

```
[ ]:
```

```
Elevation  Aspect  Slope  Horizontal_Distance_To_Hydrology  \
0      2596     51     3                      258
1      2590     56     2                      212
2      2804    139     9                      268
3      2785    155    18                      242
4      2595     45     2                      153

Vertical_Distance_To_Hydrology  Horizontal_Distance_To_Roadways  \
0                             0                             510
1                             -6                             390
2                             65                             3180
3                             118                            3090
4                             -1                             391

Hillshade_9am  Hillshade_Noon  Hillshade_3pm  \
0             221             232             148
1             220             235             151
```

2	234	238	135
3	238	238	122
4	220	234	150

	Horizontal_Distance_To_Fire_Points	...	Soil_Type32	Soil_Type33	\
0	6279	...	0	0	
1	6225	...	0	0	
2	6121	...	0	0	
3	6211	...	0	0	
4	6172	...	0	0	

	Soil_Type34	Soil_Type35	Soil_Type36	Soil_Type37	Soil_Type38	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	0	

	Soil_Type39	Soil_Type40	Cover_Type
0	0	0	5
1	0	0	5
2	0	0	2
3	0	0	2
4	0	0	5

[5 rows x 55 columns]

Explore and Data Clean

```
[ ]: df.info()
      print(df.isnull().sum())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 581012 entries, 0 to 581011
Data columns (total 55 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Elevation                                581012 non-null  int64
1   Aspect                                  581012 non-null  int64
2   Slope                                   581012 non-null  int64
3   Horizontal_Distance_To_Hydrology         581012 non-null  int64
4   Vertical_Distance_To_Hydrology           581012 non-null  int64
5   Horizontal_Distance_To_Roadways          581012 non-null  int64
6   Hillshade_9am                            581012 non-null  int64
7   Hillshade_Noon                           581012 non-null  int64
8   Hillshade_3pm                            581012 non-null  int64
9   Horizontal_Distance_To_Fire_Points       581012 non-null  int64
10  Wilderness_Area1                         581012 non-null  int64
```



11	Wilderness_Area2	581012	non-null	int64
12	Wilderness_Area3	581012	non-null	int64
13	Wilderness_Area4	581012	non-null	int64
14	Soil_Type1	581012	non-null	int64
15	Soil_Type2	581012	non-null	int64
16	Soil_Type3	581012	non-null	int64
17	Soil_Type4	581012	non-null	int64
18	Soil_Type5	581012	non-null	int64
19	Soil_Type6	581012	non-null	int64
20	Soil_Type7	581012	non-null	int64
21	Soil_Type8	581012	non-null	int64
22	Soil_Type9	581012	non-null	int64
23	Soil_Type10	581012	non-null	int64
24	Soil_Type11	581012	non-null	int64
25	Soil_Type12	581012	non-null	int64
26	Soil_Type13	581012	non-null	int64
27	Soil_Type14	581012	non-null	int64
28	Soil_Type15	581012	non-null	int64
29	Soil_Type16	581012	non-null	int64
30	Soil_Type17	581012	non-null	int64
31	Soil_Type18	581012	non-null	int64
32	Soil_Type19	581012	non-null	int64
33	Soil_Type20	581012	non-null	int64
34	Soil_Type21	581012	non-null	int64
35	Soil_Type22	581012	non-null	int64
36	Soil_Type23	581012	non-null	int64
37	Soil_Type24	581012	non-null	int64
38	Soil_Type25	581012	non-null	int64
39	Soil_Type26	581012	non-null	int64
40	Soil_Type27	581012	non-null	int64
41	Soil_Type28	581012	non-null	int64
42	Soil_Type29	581012	non-null	int64
43	Soil_Type30	581012	non-null	int64
44	Soil_Type31	581012	non-null	int64
45	Soil_Type32	581012	non-null	int64
46	Soil_Type33	581012	non-null	int64
47	Soil_Type34	581012	non-null	int64
48	Soil_Type35	581012	non-null	int64
49	Soil_Type36	581012	non-null	int64
50	Soil_Type37	581012	non-null	int64
51	Soil_Type38	581012	non-null	int64
52	Soil_Type39	581012	non-null	int64
53	Soil_Type40	581012	non-null	int64
54	Cover_Type	581012	non-null	int64

dtypes: int64(55)

memory usage: 243.8 MB

Elevation 0

Aspect 0

Slope	0
Horizontal_Distance_To_Hydrology	0
Vertical_Distance_To_Hydrology	0
Horizontal_Distance_To_Roadways	0
Hillshade_9am	0
Hillshade_Noon	0
Hillshade_3pm	0
Horizontal_Distance_To_Fire_Points	0
Wilderness_Area1	0
Wilderness_Area2	0
Wilderness_Area3	0
Wilderness_Area4	0
Soil_Type1	0
Soil_Type2	0
Soil_Type3	0
Soil_Type4	0
Soil_Type5	0
Soil_Type6	0
Soil_Type7	0
Soil_Type8	0
Soil_Type9	0
Soil_Type10	0
Soil_Type11	0
Soil_Type12	0
Soil_Type13	0
Soil_Type14	0
Soil_Type15	0
Soil_Type16	0
Soil_Type17	0
Soil_Type18	0
Soil_Type19	0
Soil_Type20	0
Soil_Type21	0
Soil_Type22	0
Soil_Type23	0
Soil_Type24	0
Soil_Type25	0
Soil_Type26	0
Soil_Type27	0
Soil_Type28	0
Soil_Type29	0
Soil_Type30	0
Soil_Type31	0
Soil_Type32	0
Soil_Type33	0
Soil_Type34	0
Soil_Type35	0
Soil_Type36	0

```

Soil_Type37          0
Soil_Type38          0
Soil_Type39          0
Soil_Type40          0
Cover_Type           0
dtype: int64

```

Feature and Target split

```

[ ]: X = df.drop('Cover_Type', axis=1)
     y = df['Cover_Type'] - 1 # Subtract 1 to make labels 0-indexed

```

Scale Numeric Feature

```

[ ]: scaler = StandardScaler()
     X_scaled = scaler.fit_transform(X)

```

Train Test Split

```

[ ]: X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
    ↪random_state=42, stratify=y)

```

Train Random Forest Classifier

```

[ ]: rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
     rf_model.fit(X_train, y_train)
     rf_preds = rf_model.predict(X_test)

```

Train XGBoost Classifier

```

[ ]: xgb_model = XGBClassifier(use_label_encoder=False, eval_metric='mlogloss',
    ↪random_state=42)
     xgb_model.fit(X_train, y_train)
     xgb_preds = xgb_model.predict(X_test)

```

```

/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[13:40:53] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

```

```
bst.update(dtrain, iteration=i, fobj=obj)
```

Evaluate Models

```

[ ]: def evaluate_model(name, y_true, y_pred):
     print(f"\n Classification Report for {name}:\n")
     print(classification_report(y_true, y_pred))
     cm = confusion_matrix(y_true, y_pred)
     disp = ConfusionMatrixDisplay(confusion_matrix=cm)
     disp.plot(cmap='Blues')
     plt.title(f'Confusion Matrix: {name}')
     plt.show()

```


```

evaluate_model("Random Forest", y_test, rf_preds)
evaluate_model("XGBoost", y_test, xgb_preds)

```

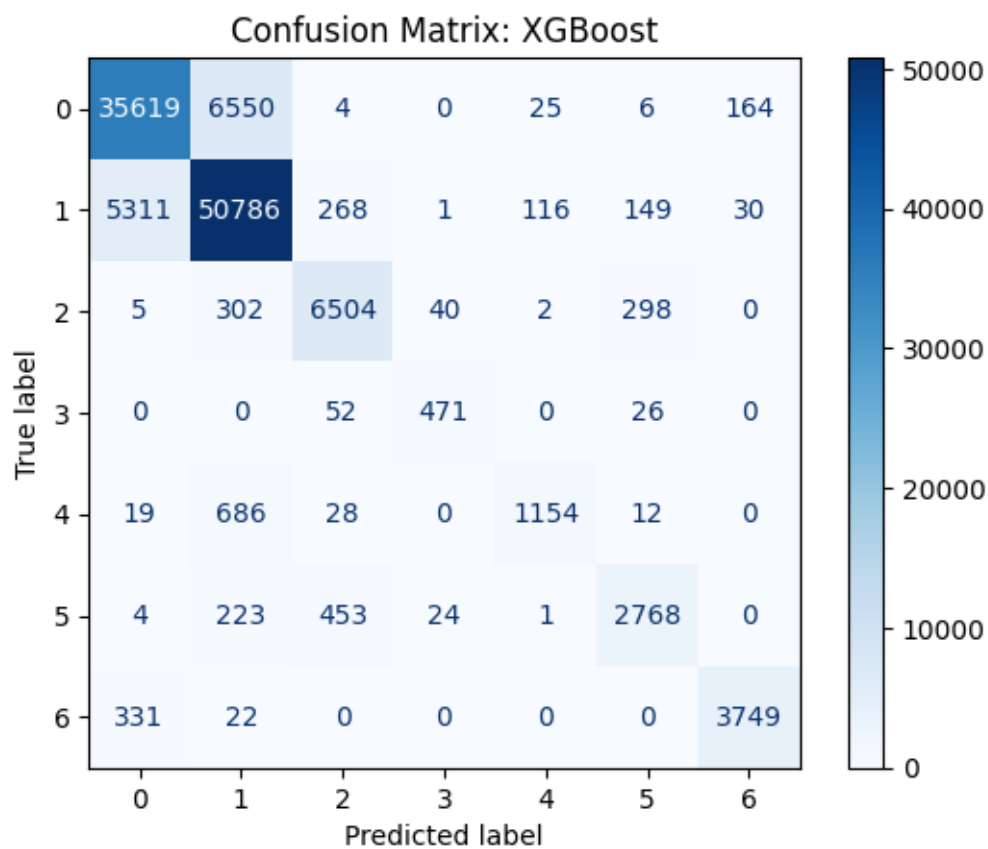
Classification Report for Random Forest:

	precision	recall	f1-score	support
0	0.96	0.94	0.95	42368
1	0.95	0.97	0.96	56661
2	0.94	0.96	0.95	7151
3	0.92	0.86	0.89	549
4	0.95	0.77	0.85	1899
5	0.93	0.89	0.91	3473
6	0.97	0.95	0.96	4102
accuracy			0.95	116203
macro avg	0.95	0.91	0.92	116203
weighted avg	0.95	0.95	0.95	116203



# Classification Report for XGBoost:

	precision	recall	f1-score	support
0	0.86	0.84	0.85	42368
1	0.87	0.90	0.88	56661
2	0.89	0.91	0.90	7151
3	0.88	0.86	0.87	549
4	0.89	0.61	0.72	1899
5	0.85	0.80	0.82	3473
6	0.95	0.91	0.93	4102
accuracy			0.87	116203
macro avg	0.88	0.83	0.85	116203
weighted avg	0.87	0.87	0.87	116203

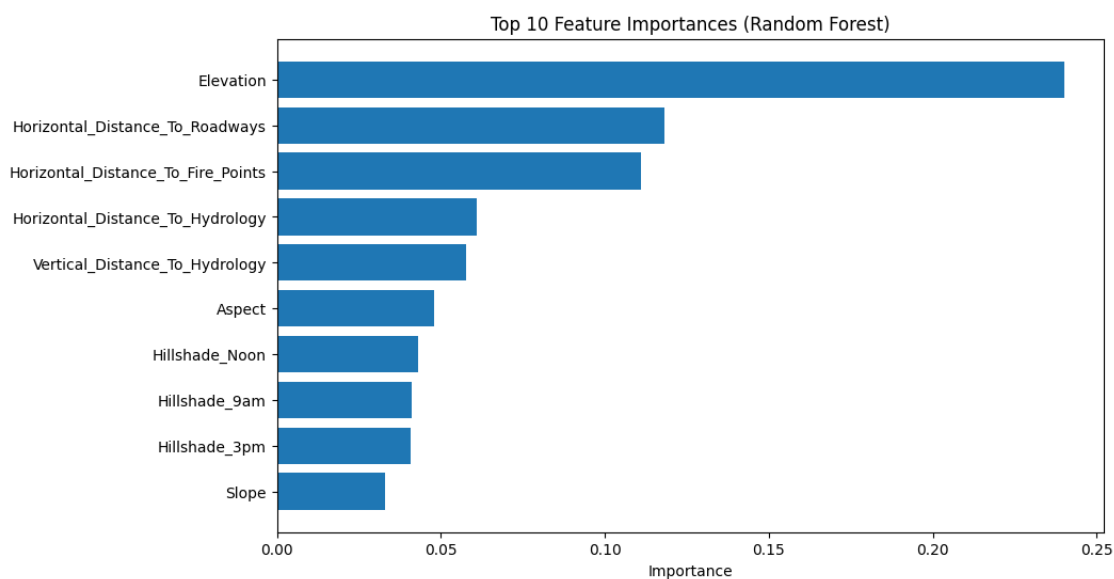


## Feature Importance

```
[ ]: # Random Forest
importances_rf = rf_model.feature_importances_
indices_rf = np.argsort(importances_rf)[-10:] # top 10

plt.figure(figsize=(10, 6))
plt.barh(range(len(indices_rf)), importances_rf[indices_rf], align='center')
plt.yticks(range(len(indices_rf)), [X.columns[i] for i in indices_rf])
plt.title("Top 10 Feature Importances (Random Forest)")
plt.xlabel("Importance")
plt.show()

# XGBoost
xgb_model.feature_importances_[:10]
```



```
[ ]: array([0.09264691, 0.0071369 , 0.00419376, 0.01318378, 0.00703175,
          0.01357988, 0.00856525, 0.01079951, 0.00502013, 0.01228627],
          dtype=float32)
```

Compare Random Forest vs. XGBoost

```
[ ]: from sklearn.metrics import accuracy_score

acc_rf = accuracy_score(y_test, rf_preds)
acc_xgb = accuracy_score(y_test, xgb_preds)

print(f" Random Forest Accuracy: {acc_rf:.4f}")
print(f" XGBoost Accuracy: {acc_xgb:.4f}")
```

Random Forest Accuracy: 0.9532

XGBoost Accuracy: 0.8696

Hyperparameter Tuning

```
[ ]: param_grid = {
    'n_estimators': [50, 100],
    'max_depth': [10, 20],
    'min_samples_split': [2, 5]
}

grid_rf = GridSearchCV(RandomForestClassifier(random_state=42), param_grid,
    cv=3, scoring='accuracy', verbose=1)
grid_rf.fit(X_train, y_train)


print(" Best Random Forest Params:", grid_rf.best_params_)
best_rf = grid_rf.best_estimator_
evaluate_model("Tuned Random Forest", y_test, best_rf.predict(X_test))
```

Fitting 3 folds for each of 8 candidates, totalling 24 fits

Best Random Forest Params: {'max\_depth': 20, 'min\_samples\_split': 2, 'n\_estimators': 100}

Classification Report for Tuned Random Forest:

	precision	recall	f1-score	support
0	0.91	0.85	0.88	42368
1	0.87	0.94	0.90	56661
2	0.92	0.93	0.92	7151
3	0.90	0.84	0.87	549
4	0.97	0.39	0.56	1899
5	0.91	0.79	0.85	3473
6	0.98	0.87	0.92	4102
accuracy			0.89	116203
macro avg	0.92	0.80	0.84	116203
weighted avg	0.89	0.89	0.89	116203



### Hyperparameter Tuning for XGBoost

```
[ ]: param_grid_xgb = {
    'n_estimators': [50, 100],
    'max_depth': [3, 6],
    'learning_rate': [0.01, 0.1]
}

grid_xgb = GridSearchCV(XGBClassifier(use_label_encoder=False,
    ↪eval_metric='mlogloss', random_state=42), param_grid_xgb, cv=3,
    ↪scoring='accuracy', verbose=1)
grid_xgb.fit(X_train, y_train)

print(" Best XGBoost Params:", grid_xgb.best_params_)
best_xgb = grid_xgb.best_estimator_
evaluate_model("Tuned XGBoost", y_test, best_xgb.predict(X_test))
```

Fitting 3 folds for each of 8 candidates, totalling 24 fits

```
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:04:00] WARNING: /workspace/src/learner.cc:738:
```



Parameters: { "use\_label\_encoder" } are not used.

```
bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:04:13] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:04:25] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:04:38] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:05:02] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:05:28] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:05:52] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:06:09] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:06:26] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:06:43] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
```

```
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:07:16] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
    bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:07:48] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
    bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:08:20] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
    bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:08:34] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
    bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:08:47] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
    bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:09:00] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
    bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:09:28] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
    bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:09:53] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
    bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:10:19] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
    bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:10:35] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:10:52] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:11:09] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:11:40] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:12:13] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

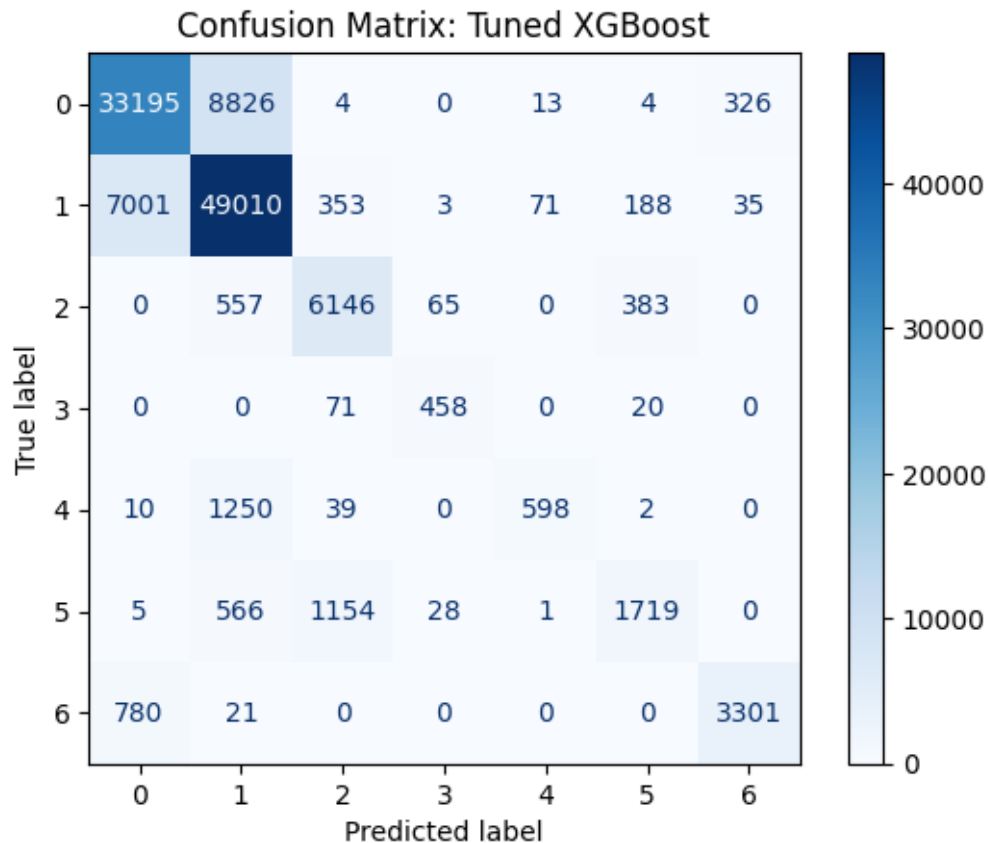
```
bst.update(dtrain, iteration=i, fobj=obj)
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[14:12:46] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)

Best XGBoost Params: {'learning_rate': 0.1, 'max_depth': 6, 'n_estimators':
100}
```

#### Classification Report for Tuned XGBoost:

	precision	recall	f1-score	support
0	0.81	0.78	0.80	42368
1	0.81	0.86	0.84	56661
2	0.79	0.86	0.82	7151
3	0.83	0.83	0.83	549
4	0.88	0.31	0.46	1899
5	0.74	0.49	0.59	3473
6	0.90	0.80	0.85	4102
accuracy			0.81	116203
macro avg	0.82	0.71	0.74	116203
weighted avg	0.81	0.81	0.81	116203



## 4 Task 4: Loan Approval Prediction Description

```
[ ]: #Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.metrics import classification_report, confusion_matrix, \
    ConfusionMatrixDisplay
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from imblearn.over_sampling import SMOTE
```

```
[ ]: #Load Dataset
import zipfile
```

```

with zipfile.ZipFile('/content/loan_approval.zip', 'r') as zip_ref:
    zip_ref.extractall('/content/')

df_train = pd.read_csv('/content/train_u6lujuX_CVtuZ9i.csv')

print("Training data:")
display(df_train.head())

```

Training data:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y

Exploratory Data Analysis(EDA)

```

[ ]: print(df_train.info())
print(df_train.isnull().sum())
print(df_train['Loan_Status'].value_counts())

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID               614 non-null   object
1   Gender                601 non-null   object
2   Married               611 non-null   object
3   Dependents            599 non-null   object
4   Education              614 non-null   object
5   Self_Employed         582 non-null   object
6   ApplicantIncome       614 non-null   int64

```

```

7   CoapplicantIncome  614 non-null   float64
8   LoanAmount         592 non-null   float64
9   Loan_Amount_Term   600 non-null   float64
10  Credit_History     564 non-null   float64
11  Property_Area      614 non-null   object
12  Loan_Status        614 non-null   object
dtypes: float64(4), int64(1), object(8)
memory usage: 62.5+ KB
None
Loan_ID          0
Gender           13
Married          3
Dependents       15
Education        0
Self_Employed    32
ApplicantIncome  0
CoapplicantIncome 0
LoanAmount       22
Loan_Amount_Term 14
Credit_History   50
Property_Area    0
Loan_Status      0
dtype: int64
Loan_Status
Y    422
N    192
Name: count, dtype: int64

```

```


[ ]: #Visualize Class Imbalance
sns.countplot(data=df_train, x='Loan_Status')
plt.title("Loan Status Distribution")

```

```

[ ]: Text(0.5, 1.0, 'Loan Status Distribution')

```



Handle Missing values

```
[ ]: # Fill numerical with median, categorical with mode
df_train['Gender'] = df_train['Gender'].fillna(df_train['Gender'].mode()[0])
df_train['Married'] = df_train['Married'].fillna(df_train['Married'].mode()[0])
df_train['Self_Employed'] = df_train['Self_Employed'].
    ↪fillna(df_train['Self_Employed'].mode()[0])
df_train['Credit_History'] = df_train['Credit_History'].
    ↪fillna(df_train['Credit_History'].mode()[0])
df_train['Loan_Amount_Term'] = df_train['Loan_Amount_Term'].
    ↪fillna(df_train['Loan_Amount_Term'].mode()[0])
df_train['LoanAmount'] = df_train['LoanAmount'].fillna(df_train['LoanAmount'].
    ↪median())

# Handle Dependents missing values with mode
df_train['Dependents'] = df_train['Dependents'].astype(str).
    ↪fillna(df_train['Dependents'].mode()[0])
```

Encode Categorical Variables

```
[ ]: le = LabelEncoder()

for col in ['Gender', 'Married', 'Education', 'Self_Employed', 'Property_Area', 'Dependents']:
    # Combine train and test data for fitting the encoder to handle unseen labels
    combined_data = pd.concat([df_train[col], df_test[col]], axis=0).
    astype(str).unique()
    le.fit(combined_data)
    df_train[col] = le.transform(df_train[col].astype(str))
    df_test[col] = le.transform(df_test[col].astype(str))

# Encode Loan_Status in train data separately as it's the target variable
df_train['Loan_Status'] = le.fit_transform(df_train['Loan_Status'])
```

Feature and Target Split

```
[ ]: X = df_train.drop(columns=['Loan_ID', 'Loan_Status'])
y = df_train['Loan_Status']
```

Train Test Split

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42, stratify=y)
```

Handle Imbalanced Data using SMOTE

```
[ ]: smote = SMOTE(random_state=42)
X_train_bal, y_train_bal = smote.fit_resample(X_train, y_train)

print("Class Distribution After SMOTE:\n", pd.Series(y_train_bal).
    value_counts())
```

Class Distribution After SMOTE:

```
Loan_Status
1    337
0    337
Name: count, dtype: int64
```

**Train Classification Model**

Logistic Regression

```
[ ]: lr = LogisticRegression(max_iter=1000)
lr.fit(X_train_bal, y_train_bal)
lr_preds = lr.predict(X_test)
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465:
ConvergenceWarning: lbfgs failed to converge (status=1):
```



STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

Decision Tree

```
[ ]: dt = DecisionTreeClassifier(random_state=42)
dt.fit(X_train_bal, y_train_bal)
dt_preds = dt.predict(X_test)
```

Evaluate Model Performance


```
[ ]: def evaluate_model(name, y_true, y_pred):
    print(f"\n {name} Classification Report:")
    print(classification_report(y_true, y_pred, target_names=['Not Approved',
    ↪ 'Approved']))

    cm = confusion_matrix(y_true, y_pred)
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['Not_
    ↪ Approved', 'Approved'])
    disp.plot(cmap='Blues')
    plt.title(f'{name} - Confusion Matrix')
    plt.show()

evaluate_model("Logistic Regression", y_test, lr_preds)
evaluate_model("Decision Tree", y_test, dt_preds)
```


Logistic Regression Classification Report:

	precision	recall	f1-score	support
Not Approved	0.65	0.68	0.67	38
Approved	0.86	0.84	0.85	85
accuracy			0.79	123
macro avg	0.75	0.76	0.76	123
weighted avg	0.79	0.79	0.79	123



#### Decision Tree Classification Report:

	precision	recall	f1-score	support
Not Approved	0.51	0.68	0.58	38
Approved	0.83	0.71	0.76	85
accuracy			0.70	123
macro avg	0.67	0.70	0.67	123
weighted avg	0.73	0.70	0.71	123



## 5 Task 5: Movie Recommendation System Description

```
[ ]: #Libraries
import pandas as pd
import numpy as np
from sklearn.metrics.pairwise import cosine_similarity
from sklearn.metrics import precision_score
from sklearn.metrics import pairwise_distances
from sklearn.decomposition import TruncatedSVD
```

Load and Merge Data

```
[ ]: #Load Dataset
import zipfile

# Extract the zip file
with zipfile.ZipFile('/content/Movie.zip', 'r') as zip_ref:
    zip_ref.extractall('/content/')
```

```

# Load ratings data
ratings_columns = ['user_id', 'item_id', 'rating', 'timestamp']
ratings_df = pd.read_csv('/content/ml-100k/u.data', sep='\t',
    ↪names=ratings_columns)

# Load movie titles data
movies_columns = ['item_id', 'title', 'release_date', 'video_release_date',
    ↪'IMDb_URL', 'unknown', 'Action', 'Adventure', 'Animation', "Children's",
    ↪'Comedy', 'Crime', 'Documentary', 'Drama', 'Fantasy', 'Film-Noir', 'Horror',
    ↪'Musical', 'Mystery', 'Romance', 'Sci-Fi', 'Thriller', 'War', 'Western']
movies_df = pd.read_csv('/content/ml-100k/u.item', sep='|',
    ↪names=movies_columns, encoding='latin-1')

print("Ratings data:")
display(ratings_df.head())

print("\nMovie titles data:")
display(movies_df.head())

```

Ratings data:

	user_id	item_id	rating	timestamp
0	196	242	3	881250949
1	186	302	3	891717742
2	22	377	1	878887116
3	244	51	2	880606923
4	166	346	1	886397596

Movie titles data:

	item_id	title	release_date	video_release_date	\
0	1	Toy Story (1995)	01-Jan-1995	NaN	
1	2	GoldenEye (1995)	01-Jan-1995	NaN	
2	3	Four Rooms (1995)	01-Jan-1995	NaN	
3	4	Get Shorty (1995)	01-Jan-1995	NaN	
4	5	Copycat (1995)	01-Jan-1995	NaN	

	IMDb_URL	unknown	Action	\
0	http://us.imdb.com/M/title-exact?Toy%20Story%2...	0	0	
1	http://us.imdb.com/M/title-exact?GoldenEye%20(...	0	1	
2	http://us.imdb.com/M/title-exact?Four%20Rooms%...	0	0	
3	http://us.imdb.com/M/title-exact?Get%20Shorty%...	0	1	
4	http://us.imdb.com/M/title-exact?Copycat%20(1995)	0	0	

	Adventure	Animation	Children's	...	Fantasy	Film-Noir	Horror	Musical	\
0	0	1	1	...	0	0	0	0	
1	1	0	0	...	0	0	0	0	
2	0	0	0	...	0	0	0	0	

3	0	0	0 ...	0	0	0	0
4	0	0	0 ...	0	0	0	0

	Mystery	Romance	Sci-Fi	Thriller	War	Western
0	0	0	0	0	0	0
1	0	0	0	1	0	0
2	0	0	0	1	0	0
3	0	0	0	0	0	0
4	0	0	0	1	0	0

[5 rows x 24 columns]

Create a User-Item Matrix

```
[ ]: # Merge ratings and movies dataframes to get movie titles
merged_df = pd.merge(ratings_df, movies_df[['item_id', 'title']], on='item_id')

# Create a User-Item Matrix
user_item_matrix = merged_df.pivot_table(index='user_id', columns='title',
    values='rating')
user_item_matrix.fillna(0, inplace=True)
display(user_item_matrix.head())
```

title	'Til There Was You (1997)	1-900 (1994)	101 Dalmatians (1996)	\
user_id				
1		0.0	0.0	2.0
2		0.0	0.0	0.0
3		0.0	0.0	0.0
4		0.0	0.0	0.0
5		0.0	0.0	2.0

title	12 Angry Men (1957)	187 (1997)	2 Days in the Valley (1996)	\
user_id				
1	5.0	0.0		0.0
2	0.0	0.0		0.0
3	0.0	2.0		0.0
4	0.0	0.0		0.0
5	0.0	0.0		0.0

title	20,000 Leagues Under the Sea (1954)	2001: A Space Odyssey (1968)	\
user_id			
1		3.0	4.0
2		0.0	0.0
3		0.0	0.0
4		0.0	0.0
5		0.0	4.0

title	3 Ninjas: High Noon At Mega Mountain (1998)	39 Steps, The (1935)	\
-------	---	----------------------	---

user_id		
1	0.0	0.0
2	1.0	0.0
3	0.0	0.0
4	0.0	0.0
5	0.0	0.0

title	...	Yankee Zulu (1994)	Year of the Horse (1997)	\
user_id	...			
1	...	0.0	0.0	
2	...	0.0	0.0	
3	...	0.0	0.0	
4	...	0.0	0.0	
5	...	0.0	0.0	

title	You So Crazy (1994)	Young Frankenstein (1974)	Young Guns (1988)	\
user_id				
1	0.0	5.0	3.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	
5	0.0	4.0	0.0	

title	Young Guns II (1990)	Young Poisoner's Handbook, The (1995)	\
user_id			
1	0.0	0.0	
2	0.0	0.0	
3	0.0	0.0	
4	0.0	0.0	
5	0.0	0.0	

title	Zeus and Roxanne (1997)	unknown	Á köldum klaka (Cold Fever) (1994)
user_id			
1	0.0	4.0	0.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0
5	0.0	4.0	0.0

[5 rows x 1664 columns]

Compute User Similarity (Cosine)

```
[ ]: from sklearn.metrics.pairwise import cosine_similarity

# Compute cosine similarity between users
user_similarity = cosine_similarity(user_item_matrix)
```

```

user_similarity_df = pd.DataFrame(user_similarity, index=user_item_matrix.
    ↪index, columns=user_item_matrix.index)

user_similarity_df.head()

```

```

[ ]: user_id      1      2      3      4      5      6      7  \
user_id
1      1.000000  0.168937  0.048388  0.064561  0.379670  0.429682  0.443097
2      0.168937  1.000000  0.113393  0.179694  0.073623  0.242106  0.108604
3      0.048388  0.113393  1.000000  0.349781  0.021592  0.074018  0.067423
4      0.064561  0.179694  0.349781  1.000000  0.031804  0.068431  0.091507
5      0.379670  0.073623  0.021592  0.031804  1.000000  0.238636  0.374733

user_id      8      9     10  ...     934     935     936  \
user_id      ...
1      0.320079  0.078385  0.377733  ...  0.372213  0.119860  0.269860
2      0.104257  0.162470  0.161273  ...  0.147095  0.310661  0.363328
3      0.084419  0.062039  0.066217  ...  0.033885  0.043453  0.167140
4      0.188060  0.101284  0.060859  ...  0.054615  0.036784  0.133619
5      0.248930  0.056847  0.201427  ...  0.340183  0.080580  0.095284

user_id     937     938     939     940     941     942     943
user_id
1      0.193343  0.197949  0.118722  0.315064  0.149086  0.181612  0.399432
2      0.410725  0.322713  0.231096  0.228793  0.162911  0.175273  0.106732
3      0.071288  0.126278  0.026758  0.164539  0.102899  0.136757  0.026990
4      0.196561  0.146058  0.030202  0.196858  0.152041  0.171538  0.058752
5      0.081053  0.148607  0.071612  0.239955  0.139595  0.153799  0.313941

[5 rows x 943 columns]

```

Recommend Movies for a Target User

```

[ ]: def recommend_movies(user_id, user_item_matrix, similarity_matrix,
    ↪n_recommendations=5):
    similar_users = similarity_matrix[user_id].sort_values(ascending=False)[1:
    ↪11] # Top 10 similar users
    weighted_ratings = pd.Series(dtype=float)

    for sim_user, similarity_score in similar_users.items():
        sim_user_ratings = user_item_matrix.loc[sim_user]
        weighted_ratings = weighted_ratings.add(sim_user_ratings *
    ↪similarity_score, fill_value=0)

    # Remove movies already rated by the user
    seen_movies = user_item_matrix.loc[user_id][user_item_matrix.loc[user_id] >
    ↪0].index

```

```
weighted_ratings = weighted_ratings.drop(seen_movies, errors='ignore')

return weighted_ratings.sort_values(ascending=False).head(n_recommendations)
```

```
[ ]: #Recommend for user id 50
recommendations = recommend_movies(50, user_item_matrix, user_similarity_df)
print("Top recommended movies for user 50:")
print(recommendations)
```

Top recommended movies for user 50:

```
title
Star Wars (1977)          9.834995
Welcome to the Dollhouse (1995)  8.124998
Postino, Il (1994)        7.531862
Bound (1996)              7.506599
Godfather, The (1972)     7.344310
dtype: float64
```

Evaluation

```
[ ]: def precision_at_k(user_id, recommended_movies, user_item_matrix, k=5):
    actual Rated = user_item_matrix.loc[user_id]
    relevant_items = actual Rated[actual Rated >= 4].index # consider rating
    ↪ >= 4 as relevant
    recommended_set = set(recommended_movies.index[:k])
    relevant_set = set(relevant_items)

    true_positives = recommended_set.intersection(relevant_set)
    return len(true_positives) / k
```

Item-Based Collaborative Filtering

```
[ ]: # Transpose the user-item matrix
item_user_matrix = user_item_matrix.T

# Compute item similarity
item_similarity = cosine_similarity(item_user_matrix)
item_similarity_df = pd.DataFrame(item_similarity, index=item_user_matrix.
    ↪ index, columns=item_user_matrix.index)

item_similarity_df.head()
```

```
[ ]: title          'Til There Was You (1997)  1-900 (1994)  \
title
'Til There Was You (1997)          1.000000    0.000000
1-900 (1994)                      0.000000    1.000000
101 Dalmatians (1996)             0.024561    0.014139
12 Angry Men (1957)              0.099561    0.009294
```



187 (1997)	0.185236	0.007354
------------	----------	----------

title	101 Dalmatians (1996)	12 Angry Men (1957)	\
title			
'Til There Was You (1997)	0.024561	0.099561	
1-900 (1994)	0.014139	0.009294	
101 Dalmatians (1996)	1.000000	0.167006	
12 Angry Men (1957)	0.167006	1.000000	
187 (1997)	0.061105	0.056822	

title	187 (1997)	2 Days in the Valley (1996)	\
title			
'Til There Was You (1997)	0.185236	0.159265	
1-900 (1994)	0.007354	0.004702	
101 Dalmatians (1996)	0.061105	0.143878	
12 Angry Men (1957)	0.056822	0.167235	
187 (1997)	1.000000	0.132327	

title	20,000 Leagues Under the Sea (1954)	\
title		
'Til There Was You (1997)	0.000000	
1-900 (1994)	0.010055	
101 Dalmatians (1996)	0.203781	
12 Angry Men (1957)	0.304078	
187 (1997)	0.042928	

title	2001: A Space Odyssey (1968)	\
title		
'Til There Was You (1997)	0.052203	
1-900 (1994)	0.067038	
101 Dalmatians (1996)	0.225803	
12 Angry Men (1957)	0.422506	
187 (1997)	0.065060	

title	3 Ninjas: High Noon At Mega Mountain (1998)	\
title		
'Til There Was You (1997)	0.000000	
1-900 (1994)	0.000000	
101 Dalmatians (1996)	0.027642	
12 Angry Men (1957)	0.072682	
187 (1997)	0.043133	

title	39 Steps, The (1935)	...	Yankee Zulu (1994)	\
title		...		
'Til There Was You (1997)	0.033326	...	0.000000	
1-900 (1994)	0.000000	...	0.152499	
101 Dalmatians (1996)	0.092337	...	0.000000	

12 Angry Men (1957)	0.394854 ...	0.060946
187 (1997)	0.027300 ...	0.000000
title Year of the Horse (1997) You So Crazy (1994) \		
title		
'Til There Was You (1997)	0.000000	0.000000
1-900 (1994)	0.015484	0.000000
101 Dalmatians (1996)	0.021965	0.030905
12 Angry Men (1957)	0.016502	0.000000
187 (1997)	0.141997	0.000000
title Young Frankenstein (1974) Young Guns (1988) \		
title		
'Til There Was You (1997)	0.027774	0.118840
1-900 (1994)	0.069284	0.018243
101 Dalmatians (1996)	0.274877	0.204267
12 Angry Men (1957)	0.403270	0.259436
187 (1997)	0.068257	0.067786
title Young Guns II (1990) \		
title		
'Til There Was You (1997)	0.142315	
1-900 (1994)	0.023408	
101 Dalmatians (1996)	0.101199	
12 Angry Men (1957)	0.145519	
187 (1997)	0.091293	
title Young Poisoner's Handbook, The (1995) \		
title		
'Til There Was You (1997)		0.029070
1-900 (1994)		0.006694
101 Dalmatians (1996)		0.056976
12 Angry Men (1957)		0.105226
187 (1997)		0.099490
title Zeus and Roxanne (1997) unknown \		
title		
'Til There Was You (1997)	0.000000	0.110208
1-900 (1994)	0.079640	0.042295
101 Dalmatians (1996)	0.172155	0.045714
12 Angry Men (1957)	0.038901	0.060101
187 (1997)	0.025184	0.142667
title Á köldum klaka (Cold Fever) (1994)		
title		
'Til There Was You (1997)		0.000000
1-900 (1994)		0.000000

101 Dalmatians (1996)	0.000000
12 Angry Men (1957)	0.081261
187 (1997)	0.096449

[5 rows x 1664 columns]

```
[ ]: #To recommend based on similar items:
def recommend_similar_items(movie_name, item_user_matrix, item_similarity_df,
    ↪n=5):
    similar_scores = item_similarity_df[movie_name].
    ↪sort_values(ascending=False)[1:n+1]
    return similar_scores
```

```
[ ]: #Example
print("Movies similar to 'Star Wars (1977)':")
print(recommend_similar_items("Star Wars (1977)", item_user_matrix,
    ↪item_similarity_df))
```

Movies similar to 'Star Wars (1977)':

title	
Return of the Jedi (1983)	0.884476
Raiders of the Lost Ark (1981)	0.764885
Empire Strikes Back, The (1980)	0.749819
Toy Story (1995)	0.734572
Godfather, The (1972)	0.697332

Name: Star Wars (1977), dtype: float64

Matrix Factorization using SVD

```
[ ]: from scipy.sparse.linalg import svds
import numpy as np

# Convert to numpy matrix
R = user_item_matrix.values
user_ratings_mean = np.mean(R, axis=1)
R_demeaned = R - user_ratings_mean.reshape(-1, 1)

# Perform SVD
U, sigma, Vt = svds(R_demeaned, k=50)
sigma = np.diag(sigma)

# Reconstruct the predicted ratings
predicted_ratings = np.dot(np.dot(U, sigma), Vt) + user_ratings_mean.
    ↪reshape(-1, 1)
predicted_df = pd.DataFrame(predicted_ratings, columns=user_item_matrix.
    ↪columns, index=user_item_matrix.index)
```

```
[ ]: #Recommend for User 50 using SVD
def recommend_from_svd(user_id, predictions_df, user_item_matrix, n=5):
    user_row = predictions_df.loc[user_id]
    watched_movies = user_item_matrix.loc[user_id][user_item_matrix.
↳loc[user_id] > 0].index
    recommendations = user_row.drop(watched_movies, errors='ignore').
↳sort_values(ascending=False).head(n)
    return recommendations

print("SVD Recommendations for User 50:")
print(recommend_from_svd(50, predicted_df, user_item_matrix))
```

SVD Recommendations for User 50:

```
title
Bound (1996)                1.306506
Twelve Monkeys (1995)      1.252929
Boogie Nights (1997)       1.173483
Welcome to the Dollhouse (1995) 1.139967
Game, The (1997)           0.979868
Name: 50, dtype: float64
```

```
[ ]:
```

## Task 6: Music Genre Classification Description

```
[ ]: import zipfile
import os

zip_path = "/content/drive/MyDrive/music_genre.zip"
extract_path = "/content/music_genre"

with zipfile.ZipFile(zip_path, 'r') as zip_ref:
    zip_ref.extractall(extract_path)

# Check extracted files
for root, dirs, files in os.walk(extract_path):
    print(root, len(files))
```

```
/content/music_genre 0
/content/music_genre/Data 2
/content/music_genre/Data/images_original 0
/content/music_genre/Data/images_original/jazz 99
/content/music_genre/Data/images_original/blues 100
/content/music_genre/Data/images_original/hiphop 100
/content/music_genre/Data/images_original/pop 100
/content/music_genre/Data/images_original/classical 100
/content/music_genre/Data/images_original/reggae 100
/content/music_genre/Data/images_original/metal 100
```

```

/content/music_genre/Data/images_original/disco 100
/content/music_genre/Data/images_original/country 100
/content/music_genre/Data/images_original/rock 100
/content/music_genre/Data/genres_original 0
/content/music_genre/Data/genres_original/jazz 100
/content/music_genre/Data/genres_original/blues 100
/content/music_genre/Data/genres_original/hiphop 100
/content/music_genre/Data/genres_original/pop 100
/content/music_genre/Data/genres_original/classical 100
/content/music_genre/Data/genres_original/reggae 100
/content/music_genre/Data/genres_original/metal 100
/content/music_genre/Data/genres_original/disco 100
/content/music_genre/Data/genres_original/country 100
/content/music_genre/Data/genres_original/rock 100

```

```

[ ]: import glob
import os

# Path to audio files
audio_path = "/content/music_genre/Data/genres_original"

# Get genre folders
genres = [g for g in os.listdir(audio_path) if os.path.isdir(os.path.
    ↪join(audio_path, g))]
print("Genres found:", genres)

filepaths = []
labels = []

for genre in genres:
    files = glob.glob(os.path.join(audio_path, genre, "*.wav"))
    for f in files:
        filepaths.append(f)
        labels.append(genre)

print(f"Found {len(filepaths)} audio files.")

```

Genres found: ['jazz', 'blues', 'hiphop', 'pop', 'classical', 'reggae', 'metal', 'disco', 'country', 'rock']

Found 1000 audio files.

Audio Feature Extraction (MFCC) – Tabular Approach

```

[ ]: import librosa
import numpy as np
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier

```

```

from sklearn.metrics import classification_report, accuracy_score
import warnings
warnings.filterwarnings('ignore')

# Extract MFCC Features
X = []
y = []

for idx, filepath in enumerate(filepaths):
    try:
        audio, sr = librosa.load(filepath, duration=30) # load 30 sec clip
        mfccs = librosa.feature.mfcc(y=audio, sr=sr, n_mfcc=20)
        mfccs_mean = np.mean(mfccs.T, axis=0) # mean over time
        X.append(mfccs_mean)
        y.append(labels[idx])
    except Exception as e:
        print(f"Error processing {filepath}: {e}")

X = np.array(X)
y = np.array(y)

print("Feature matrix shape:", X.shape)
print("Labels count:", len(y))

```

Error processing /content/music\_genre/Data/genres\_original/jazz/jazz.00054.wav:  
Feature matrix shape: (999, 20)  
Labels count: 999

```

[ ]: # Encode Labels
label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y)

# Scale Features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Train/Test Split
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y_encoded, test_size=0.2, random_state=42
)

```

Train a Classifier

```

[ ]: # Train Classifier (Random Forest)
clf = RandomForestClassifier(n_estimators=200, random_state=42)
clf.fit(X_train, y_train)

```

```

[ ]: RandomForestClassifier(n_estimators=200, random_state=42)

```

## Evaluation

```
[ ]: y_pred = clf.predict(X_test)
print("\nAccuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred,
↳target_names=label_encoder.classes_))
```

Accuracy: 0.625

### Classification Report:

	precision	recall	f1-score	support
blues	0.55	0.55	0.55	11
classical	0.80	0.80	0.80	15
country	0.46	0.50	0.48	22
disco	0.50	0.68	0.58	19
hiphop	0.62	0.54	0.58	24
jazz	0.76	0.62	0.68	21
metal	0.62	0.72	0.67	18
pop	0.76	0.83	0.79	23
reggae	0.62	0.67	0.64	27
rock	0.64	0.35	0.45	20
accuracy			0.62	200
macro avg	0.63	0.63	0.62	200
weighted avg	0.63	0.62	0.62	200

## Generate Spectrogram Images – CNN Approach

```
[ ]: import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.layers import Dense, GlobalAveragePooling2D, Dropout
from tensorflow.keras.models import Model

# Paths
img_dir = "/content/music_genre/Data/images_original"

# Data Generator with Augmentation
datagen = ImageDataGenerator(
    rescale=1./255,
    validation_split=0.2,
    rotation_range=15,
    width_shift_range=0.1,
    height_shift_range=0.1,
    horizontal_flip=True)
```

```

)

train_gen = datagen.flow_from_directory(
    img_dir,
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical',
    subset='training'
)

val_gen = datagen.flow_from_directory(
    img_dir,
    target_size=(224, 224),
    batch_size=32,
    class_mode='categorical',
    subset='validation'
)

```

Found 800 images belonging to 10 classes.  
Found 199 images belonging to 10 classes.

```

[ ]: # Load MobileNetV2 Pretrained Model
base_model = MobileNetV2(weights='imagenet', include_top=False,
    ↪input_shape=(224, 224, 3))
base_model.trainable = False # freeze base layers

```

```

[ ]: # Build Model
x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dropout(0.3)(x)
predictions = Dense(train_gen.num_classes, activation='softmax')(x)
model = Model(inputs=base_model.input, outputs=predictions)

```

```

[ ]: # Compile
model.compile(optimizer='adam', loss='categorical_crossentropy',
    ↪metrics=['accuracy'])

# Train
history = model.fit(
    train_gen,
    validation_data=val_gen,
    epochs=5
)

```

Epoch 1/5  
25/25                      34s 919ms/step -  
accuracy: 0.1312 - loss: 2.5960 - val\_accuracy: 0.2915 - val\_loss: 1.9250



Epoch 2/5  
25/25 14s 567ms/step -  
accuracy: 0.2975 - loss: 1.9927 - val\_accuracy: 0.4472 - val\_loss: 1.6598  
Epoch 3/5  
25/25 14s 565ms/step -  
accuracy: 0.3894 - loss: 1.6612 - val\_accuracy: 0.4623 - val\_loss: 1.5527  
Epoch 4/5  
25/25 15s 623ms/step -  
accuracy: 0.3993 - loss: 1.6027 - val\_accuracy: 0.4925 - val\_loss: 1.4535  
Epoch 5/5  
25/25 14s 565ms/step -  
accuracy: 0.5046 - loss: 1.4637 - val\_accuracy: 0.4975 - val\_loss: 1.4372

```
[ ]: # Evaluate  
loss, acc = model.evaluate(val_gen)  
print(f"\nValidation Accuracy: {acc:.2f}")
```

7/7 4s 513ms/step -  
accuracy: 0.5191 - loss: 1.3870

Validation Accuracy: 0.50

```
[ ]:
```

## Task 7: Sales Forecasting Description

```
[ ]: import pandas as pd  
import zipfile  
  
# Load dataset  
# Extract the zip file  
with zipfile.ZipFile('/content/wallmart.zip', 'r') as zip_ref:  
    zip_ref.extractall('/content/')  
  
# Load the training data  
df = pd.read_csv("/content/train.csv")
```

### Data Preparation

```
[ ]: # Convert date column to datetime  
df['Date'] = pd.to_datetime(df['Date'])  
  
# Sort by date for time series  
df = df.sort_values(by='Date')
```

### Feature Engineering

```
[ ]: # Time-based features
df['Month'] = df['Date'].dt.month
df['Week'] = df['Date'].dt.isocalendar().week
df['DayOfWeek'] = df['Date'].dt.dayofweek

# Lag features
df['Sales_Lag_1'] = df.groupby(['Store', 'Dept'])['Weekly_Sales'].shift(1)
df['Sales_Lag_2'] = df.groupby(['Store', 'Dept'])['Weekly_Sales'].shift(2)

# Rolling averages
df['Rolling_Mean_4'] = df.groupby(['Store', 'Dept'])['Weekly_Sales'].shift(1).
    ↪rolling(4).mean()
```

Handle Missing Values

```
[ ]: #Lag and rolling averages will create NaN for the first few rows → drop them:
df = df.dropna()
```

Train-Test Split (Time-Aware)

```
[ ]: #Use time-aware split so the future is not leaked into the past
from sklearn.model_selection import TimeSeriesSplit

# Example split
tscv = TimeSeriesSplit(n_splits=5)
```

Model 1 — Baseline Linear Regression

```
[ ]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
import numpy as np

features = ['Month', 'Week', 'DayOfWeek', 'Sales_Lag_1', 'Sales_Lag_2',
    ↪'Rolling_Mean_4']
target = 'Weekly_Sales'

X = df[features]
y = df[target]

# Use last 20% of data as test
split_point = int(len(X)*0.8)
X_train, X_test = X[:split_point], X[split_point:]
y_train, y_test = y[:split_point], y[split_point:]

model_lr = LinearRegression()
model_lr.fit(X_train, y_train)
y_pred_lr = model_lr.predict(X_test)
```

```
rmse_lr = np.sqrt(mean_squared_error(y_test, y_pred_lr))
print("Linear Regression RMSE:", rmse_lr)
```

Linear Regression RMSE: 3464.0457885246215

Model 2 — XGBoost

```
[ ]: from xgboost import XGBRegressor

model_xgb = XGBRegressor(n_estimators=200, learning_rate=0.1, max_depth=6)
model_xgb.fit(X_train, y_train)
y_pred_xgb = model_xgb.predict(X_test)


rmse_xgb = np.sqrt(mean_squared_error(y_test, y_pred_xgb))
print("XGBoost RMSE:", rmse_xgb)
```

XGBoost RMSE: 3220.954432468477

Visualization — Actual vs Predicted

```
[ ]: import matplotlib.pyplot as plt


plt.figure(figsize=(12,6))
plt.plot(y_test.values, label='Actual Sales', color='blue')
plt.plot(y_pred_xgb, label='Predicted Sales (XGBoost)', color='red')
plt.legend()
plt.title('Actual vs Predicted Sales')
plt.xlabel('Time')
plt.ylabel('Weekly Sales')
plt.show()
```



## Seasonal Decomposition

```
[ ]: from statsmodels.tsa.seasonal import seasonal_decompose

df_store = df[df['Store'] == 1]
result = seasonal_decompose(df_store['Weekly_Sales'], model='additive',
                             period=52)
result.plot()
plt.show()
```



```
[ ]:
```

## Task 8: Traffic Sign Recognition Description

```
[ ]: #Libraries
import os
import cv2
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
```

```
from tensorflow.keras.utils import to_categorical
```

Unzip the Data

```
[ ]: import zipfile

# Define the path to the zip file
zip_path = '/content/drive/MyDrive/traffic.zip'
extract_path = '/content/drive/MyDrive/traffic/'

# Unzipping the file
with zipfile.ZipFile(zip_path, 'r') as zip_ref:
    zip_ref.extractall(extract_path)

print("Dataset unzipped successfully.")
```

Dataset unzipped successfully.

```
[ ]: # List the contents of the extracted folder
os.listdir(extract_path)
```

```
[ ]: ['Meta.csv',
      'Meta',
      'Test.csv',
      'Test',
      'Train.csv',
      'Train',
      'meta',
      'test',
      'train']
```

Load the Data

```
[ ]: # Path to dataset
extract_path = '/content/drive/MyDrive/traffic/'

# Paths to CSV files
train_csv = os.path.join(extract_path, 'Train.csv')
test_csv = os.path.join(extract_path, 'Test.csv')

# Read CSVs
train_df = pd.read_csv(train_csv)
test_df = pd.read_csv(test_csv)

# Remove any extra spaces in column names (important!)
train_df.columns = train_df.columns.str.strip()
test_df.columns = test_df.columns.str.strip()
```

```
print(train_df.head())
```

	Width	Height	Roi.X1	Roi.Y1	Roi.X2	Roi.Y2	ClassId	\
0	27	26	5	5	22	20	20	
1	28	27	5	6	23	22	20	
2	29	26	6	5	24	21	20	
3	28	27	5	6	23	22	20	
4	28	26	5	5	23	21	20	

	Path
0	Train/20/00020_00000_00000.png
1	Train/20/00020_00000_00001.png
2	Train/20/00020_00000_00002.png
3	Train/20/00020_00000_00003.png
4	Train/20/00020_00000_00004.png

Load & Preprocess Images

```
[ ]: IMG_SIZE = 32

# Load training images
train_images = []
train_labels = []

for idx, row in train_df.iterrows():
    img_path = os.path.join(extract_path, row['Path'])
    img = cv2.imread(img_path)
    img = cv2.resize(img, (IMG_SIZE, IMG_SIZE))
    train_images.append(img)
    train_labels.append(row['ClassId'])

X_train = np.array(train_images) / 255.0 # normalize
y_train = to_categorical(np.array(train_labels))
```

Train-validation split

```
[ ]: X_train, X_val, y_train, y_val = train_test_split(
    X_train, y_train, test_size=0.2, random_state=42
)

print("Train shape:", X_train.shape)
print("Validation shape:", X_val.shape)
```

Train shape: (31367, 32, 32, 3)

Validation shape: (7842, 32, 32, 3)

Build & compile CNN

```
[ ]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
↳Dropout

model = Sequential([
    Conv2D(32, (3,3), activation='relu', input_shape=(IMG_SIZE, IMG_SIZE, 3)),
    MaxPooling2D(2,2),
    Conv2D(64, (3,3), activation='relu'),
    MaxPooling2D(2,2),
    Flatten(),
    Dense(128, activation='relu'),
    Dropout(0.5),
    Dense(y_train.shape[1], activation='softmax')
])

model.compile(optimizer='adam', loss='categorical_crossentropy',
↳metrics=['accuracy'])
model.summary()
```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base\_conv.py:113: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten (Flatten)	(None, 2304)	0
dense (Dense)	(None, 128)	295,040
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 43)	5,547

Total params: 319,979 (1.22 MB)

Trainable params: 319,979 (1.22 MB)

Non-trainable params: 0 (0.00 B)

Train the Model

```
[ ]: history = model.fit(  
    X_train, y_train,  
    validation_data=(X_val, y_val),  
    epochs=3,  
    batch_size=32  
)
```

Epoch 1/3

981/981 17s 9ms/step -  
accuracy: 0.3413 - loss: 2.4550 - val\_accuracy: 0.9118 - val\_loss: 0.4010

Epoch 2/3

981/981 5s 5ms/step -  
accuracy: 0.8233 - loss: 0.5659 - val\_accuracy: 0.9651 - val\_loss: 0.1538

Epoch 3/3

981/981 4s 4ms/step -  
accuracy: 0.8950 - loss: 0.3371 - val\_accuracy: 0.9776 - val\_loss: 0.0964

Test Data Evaluation

```
[ ]: test_images = []  
    test_labels = []  
  
    for idx, row in test_df.iterrows():  
        img_path = os.path.join(extract_path, row['Path'])  
        img = cv2.imread(img_path)  
        img = cv2.resize(img, (IMG_SIZE, IMG_SIZE))  
        test_images.append(img)  
        test_labels.append(row['ClassId'])  
  
    X_test = np.array(test_images) / 255.0  
    y_test = to_categorical(np.array(test_labels))  
  
    loss, acc = model.evaluate(X_test, y_test)  
    print(f"Test Accuracy: {acc*100:.2f}%")
```

395/395 2s 6ms/step -

accuracy: 0.9318 - loss: 0.2764

Test Accuracy: 93.10%



## Add Data Augmentation

```
[ ]: from tensorflow.keras.preprocessing.image import ImageDataGenerator

datagen = ImageDataGenerator(
    rotation_range=10,
    zoom_range=0.1,
    width_shift_range=0.1,
    height_shift_range=0.1,
    shear_range=0.1,
    horizontal_flip=False,
    fill_mode='nearest'
)

datagen.fit(X_train)

# Retrain the same CNN with augmentation
history_aug = model.fit(
    datagen.flow(X_train, y_train, batch_size=32),
    validation_data=(X_val, y_val),
    epochs=10
)
```

Epoch 1/10

/usr/local/lib/python3.11/dist-

packages/keras/src/trainers/data\_adapters/py\_dataset\_adapter.py:121:

UserWarning: Your `PyDataset` class should call `super().\_\_init\_\_(\*\*kwargs)` in its constructor. `\*\*kwargs` can include `workers`, `use\_multiprocessing`, `max\_queue\_size`. Do not pass these arguments to `fit()`, as they will be ignored.

self.\_warn\_if\_super\_not\_called()

981/981 26s 24ms/step -

accuracy: 0.6251 - loss: 1.3288 - val\_accuracy: 0.9652 - val\_loss: 0.1443

Epoch 2/10

981/981 40s 25ms/step -

accuracy: 0.7540 - loss: 0.7829 - val\_accuracy: 0.9759 - val\_loss: 0.1234

Epoch 3/10

981/981 29s 29ms/step -

accuracy: 0.7985 - loss: 0.6382 - val\_accuracy: 0.9753 - val\_loss: 0.1136

Epoch 4/10

981/981 23s 23ms/step -

accuracy: 0.8245 - loss: 0.5410 - val\_accuracy: 0.9790 - val\_loss: 0.0944

Epoch 5/10

981/981 38s 20ms/step -

accuracy: 0.8524 - loss: 0.4723 - val\_accuracy: 0.9856 - val\_loss: 0.0658

Epoch 6/10

981/981 24s 24ms/step -

```

accuracy: 0.8602 - loss: 0.4257 - val_accuracy: 0.9839 - val_loss: 0.0689
Epoch 7/10
981/981          22s 22ms/step -
accuracy: 0.8788 - loss: 0.3758 - val_accuracy: 0.9890 - val_loss: 0.0479
Epoch 8/10
981/981          20s 20ms/step -
accuracy: 0.8897 - loss: 0.3459 - val_accuracy: 0.9888 - val_loss: 0.0482
Epoch 9/10
981/981          21s 22ms/step -
accuracy: 0.9003 - loss: 0.3235 - val_accuracy: 0.9893 - val_loss: 0.0431
Epoch 10/10
981/981          21s 21ms/step -
accuracy: 0.9026 - loss: 0.3022 - val_accuracy: 0.9911 - val_loss: 0.0345

```

MobileNetV2 Transfer Learning

```

[ ]: from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.models import Model
from tensorflow.keras.layers import GlobalAveragePooling2D

# Load MobileNetV2 without top layers
base_model = MobileNetV2(weights='imagenet', include_top=False,
    ↪input_shape=(IMG_SIZE, IMG_SIZE, 3))
base_model.trainable = False # freeze base model

# Add custom layers on top
x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(128, activation='relu')(x)
x = Dropout(0.5)(x)
predictions = Dense(y_train.shape[1], activation='softmax')(x)

mobilenet_model = Model(inputs=base_model.input, outputs=predictions)

mobilenet_model.compile(optimizer='adam', loss='categorical_crossentropy',
    ↪metrics=['accuracy'])

mobilenet_model.summary()

```

/tmp/ipython-input-859176245.py:6: UserWarning: `input\_shape` is undefined or non-square, or `rows` is not in [96, 128, 160, 192, 224]. Weights for input shape (224, 224) will be loaded as the default.

```

base_model = MobileNetV2(weights='imagenet', include_top=False,
input_shape=(IMG_SIZE, IMG_SIZE, 3))

```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/mobilenet\\_v2/mobilenet\\_v2\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_1.0\\_224\\_no\\_top.h5](https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_224_no_top.h5)

9406464/9406464

0s

Ous/step

Model: "functional\_1"

Layer (type)	Output Shape	Param #	Connected to
input_layer_1 (InputLayer)	(None, 32, 32, 3)	0	-
Conv1 (Conv2D)	(None, 16, 16, 32)	864	input_layer_1[0]...
bn_Conv1 (BatchNormalizatio...	(None, 16, 16, 32)	128	Conv1[0][0]
Conv1_relu (ReLU)	(None, 16, 16, 32)	0	bn_Conv1[0][0]
expanded_conv_dept... (DepthwiseConv2D)	(None, 16, 16, 32)	288	Conv1_relu[0][0]
expanded_conv_dept... (BatchNormalizatio...	(None, 16, 16, 32)	128	expanded_conv_de...
expanded_conv_dept... (ReLU)	(None, 16, 16, 32)	0	expanded_conv_de...
expanded_conv_proj... (Conv2D)	(None, 16, 16, 16)	512	expanded_conv_de...
expanded_conv_proj... (BatchNormalizatio...	(None, 16, 16, 16)	64	expanded_conv_pr...
block_1_expand (Conv2D)	(None, 16, 16, 96)	1,536	expanded_conv_pr...
block_1_expand_BN (BatchNormalizatio...	(None, 16, 16, 96)	384	block_1_expand[0...
block_1_expand_relu (ReLU)	(None, 16, 16, 96)	0	block_1_expand_B...
block_1_pad (ZeroPadding2D)	(None, 17, 17, 96)	0	block_1_expand_r...
block_1_depthwise (DepthwiseConv2D)	(None, 8, 8, 96)	864	block_1_pad[0][0]

block_1_depthwise_... (BatchNormalizatio...	(None, 8, 8, 96)	384	block_1_depthwis...
block_1_depthwise_... (ReLU)	(None, 8, 8, 96)	0	block_1_depthwis...
block_1_project (Conv2D)	(None, 8, 8, 24)	2,304	block_1_depthwis...
block_1_project_BN (BatchNormalizatio...	(None, 8, 8, 24)	96	block_1_project[...
block_2_expand (Conv2D)	(None, 8, 8, 144)	3,456	block_1_project_...
block_2_expand_BN (BatchNormalizatio...	(None, 8, 8, 144)	576	block_2_expand[0...
block_2_expand_relu (ReLU)	(None, 8, 8, 144)	0	block_2_expand_B...
block_2_depthwise (DepthwiseConv2D)	(None, 8, 8, 144)	1,296	block_2_expand_r...
block_2_depthwise_... (BatchNormalizatio...	(None, 8, 8, 144)	576	block_2_depthwis...
block_2_depthwise_... (ReLU)	(None, 8, 8, 144)	0	block_2_depthwis...
block_2_project (Conv2D)	(None, 8, 8, 24)	3,456	block_2_depthwis...
block_2_project_BN (BatchNormalizatio...	(None, 8, 8, 24)	96	block_2_project[...
block_2_add (Add)	(None, 8, 8, 24)	0	block_1_project_... block_2_project_...
block_3_expand (Conv2D)	(None, 8, 8, 144)	3,456	block_2_add[0][0]
block_3_expand_BN (BatchNormalizatio...	(None, 8, 8, 144)	576	block_3_expand[0...
block_3_expand_relu (ReLU)	(None, 8, 8, 144)	0	block_3_expand_B...

block_3_pad (ZeroPadding2D)	(None, 9, 9, 144)	0	block_3_expand_r...
block_3_depthwise (DepthwiseConv2D)	(None, 4, 4, 144)	1,296	block_3_pad[0][0]
block_3_depthwise_... (BatchNormalizatio...	(None, 4, 4, 144)	576	block_3_depthwis...
block_3_depthwise_... (ReLU)	(None, 4, 4, 144)	0	block_3_depthwis...
block_3_project (Conv2D)	(None, 4, 4, 32)	4,608	block_3_depthwis...
block_3_project_BN (BatchNormalizatio...	(None, 4, 4, 32)	128	block_3_project[...
block_4_expand (Conv2D)	(None, 4, 4, 192)	6,144	block_3_project_...
block_4_expand_BN (BatchNormalizatio...	(None, 4, 4, 192)	768	block_4_expand[0...
block_4_expand_relu (ReLU)	(None, 4, 4, 192)	0	block_4_expand_B...
block_4_depthwise (DepthwiseConv2D)	(None, 4, 4, 192)	1,728	block_4_expand_r...
block_4_depthwise_... (BatchNormalizatio...	(None, 4, 4, 192)	768	block_4_depthwis...
block_4_depthwise_... (ReLU)	(None, 4, 4, 192)	0	block_4_depthwis...
block_4_project (Conv2D)	(None, 4, 4, 32)	6,144	block_4_depthwis...
block_4_project_BN (BatchNormalizatio...	(None, 4, 4, 32)	128	block_4_project[...
block_4_add (Add)	(None, 4, 4, 32)	0	block_3_project_... block_4_project_...
block_5_expand (Conv2D)	(None, 4, 4, 192)	6,144	block_4_add[0][0]

block_5_expand_BN (BatchNormalizatio...	(None, 4, 4, 192)	768	block_5_expand[0...
block_5_expand_relu (ReLU)	(None, 4, 4, 192)	0	block_5_expand_B...
block_5_depthwise (DepthwiseConv2D)	(None, 4, 4, 192)	1,728	block_5_expand_r...
block_5_depthwise_... (BatchNormalizatio...	(None, 4, 4, 192)	768	block_5_depthwis...
block_5_depthwise_... (ReLU)	(None, 4, 4, 192)	0	block_5_depthwis...
block_5_project (Conv2D)	(None, 4, 4, 32)	6,144	block_5_depthwis...
block_5_project_BN (BatchNormalizatio...	(None, 4, 4, 32)	128	block_5_project[...
block_5_add (Add)	(None, 4, 4, 32)	0	block_4_add[0][0... block_5_project_...
block_6_expand (Conv2D)	(None, 4, 4, 192)	6,144	block_5_add[0][0]
block_6_expand_BN (BatchNormalizatio...	(None, 4, 4, 192)	768	block_6_expand[0...
block_6_expand_relu (ReLU)	(None, 4, 4, 192)	0	block_6_expand_B...
block_6_pad (ZeroPadding2D)	(None, 5, 5, 192)	0	block_6_expand_r...
block_6_depthwise (DepthwiseConv2D)	(None, 2, 2, 192)	1,728	block_6_pad[0][0]
block_6_depthwise_... (BatchNormalizatio...	(None, 2, 2, 192)	768	block_6_depthwis...
block_6_depthwise_... (ReLU)	(None, 2, 2, 192)	0	block_6_depthwis...
block_6_project (Conv2D)	(None, 2, 2, 64)	12,288	block_6_depthwis...

block_6_project_BN (BatchNormalizatio...	(None, 2, 2, 64)	256	block_6_project[...
block_7_expand (Conv2D)	(None, 2, 2, 384)	24,576	block_6_project_...
block_7_expand_BN (BatchNormalizatio...	(None, 2, 2, 384)	1,536	block_7_expand[0...
block_7_expand_relu (ReLU)	(None, 2, 2, 384)	0	block_7_expand_B...
block_7_depthwise (DepthwiseConv2D)	(None, 2, 2, 384)	3,456	block_7_expand_r...
block_7_depthwise_... (BatchNormalizatio...	(None, 2, 2, 384)	1,536	block_7_depthwis...
block_7_depthwise_... (ReLU)	(None, 2, 2, 384)	0	block_7_depthwis...
block_7_project (Conv2D)	(None, 2, 2, 64)	24,576	block_7_depthwis...
block_7_project_BN (BatchNormalizatio...	(None, 2, 2, 64)	256	block_7_project[...
block_7_add (Add)	(None, 2, 2, 64)	0	block_6_project_... block_7_project_...
block_8_expand (Conv2D)	(None, 2, 2, 384)	24,576	block_7_add[0][0]
block_8_expand_BN (BatchNormalizatio...	(None, 2, 2, 384)	1,536	block_8_expand[0...
block_8_expand_relu (ReLU)	(None, 2, 2, 384)	0	block_8_expand_B...
block_8_depthwise (DepthwiseConv2D)	(None, 2, 2, 384)	3,456	block_8_expand_r...
block_8_depthwise_... (BatchNormalizatio...	(None, 2, 2, 384)	1,536	block_8_depthwis...
block_8_depthwise_... (ReLU)	(None, 2, 2, 384)	0	block_8_depthwis...

block_8_project (Conv2D)	(None, 2, 2, 64)	24,576	block_8_depthwis...
block_8_project_BN (BatchNormalizatio...	(None, 2, 2, 64)	256	block_8_project[...
block_8_add (Add)	(None, 2, 2, 64)	0	block_7_add[0][0... block_8_project_...
block_9_expand (Conv2D)	(None, 2, 2, 384)	24,576	block_8_add[0][0]
block_9_expand_BN (BatchNormalizatio...	(None, 2, 2, 384)	1,536	block_9_expand[0...
block_9_expand_relu (ReLU)	(None, 2, 2, 384)	0	block_9_expand_B...
block_9_depthwise (DepthwiseConv2D)	(None, 2, 2, 384)	3,456	block_9_expand_r...
block_9_depthwise_... (BatchNormalizatio...	(None, 2, 2, 384)	1,536	block_9_depthwis...
block_9_depthwise_... (ReLU)	(None, 2, 2, 384)	0	block_9_depthwis...
block_9_project (Conv2D)	(None, 2, 2, 64)	24,576	block_9_depthwis...
block_9_project_BN (BatchNormalizatio...	(None, 2, 2, 64)	256	block_9_project[...
block_9_add (Add)	(None, 2, 2, 64)	0	block_8_add[0][0... block_9_project_...
block_10_expand (Conv2D)	(None, 2, 2, 384)	24,576	block_9_add[0][0]
block_10_expand_BN (BatchNormalizatio...	(None, 2, 2, 384)	1,536	block_10_expand[...
block_10_expand_re... (ReLU)	(None, 2, 2, 384)	0	block_10_expand_...
block_10_depthwise (DepthwiseConv2D)	(None, 2, 2, 384)	3,456	block_10_expand_...



block_10_depthwise... (BatchNormalizatio...	(None, 2, 2, 384)	1,536	block_10_depthwi...
block_10_depthwise... (ReLU)	(None, 2, 2, 384)	0	block_10_depthwi...
block_10_project (Conv2D)	(None, 2, 2, 96)	36,864	block_10_depthwi...
block_10_project_BN (BatchNormalizatio...	(None, 2, 2, 96)	384	block_10_project...
block_11_expand (Conv2D)	(None, 2, 2, 576)	55,296	block_10_project...
block_11_expand_BN (BatchNormalizatio...	(None, 2, 2, 576)	2,304	block_11_expand[...
block_11_expand_re... (ReLU)	(None, 2, 2, 576)	0	block_11_expand_...
block_11_depthwise (DepthwiseConv2D)	(None, 2, 2, 576)	5,184	block_11_expand_...
block_11_depthwise... (BatchNormalizatio...	(None, 2, 2, 576)	2,304	block_11_depthwi...
block_11_depthwise... (ReLU)	(None, 2, 2, 576)	0	block_11_depthwi...
block_11_project (Conv2D)	(None, 2, 2, 96)	55,296	block_11_depthwi...
block_11_project_BN (BatchNormalizatio...	(None, 2, 2, 96)	384	block_11_project...
block_11_add (Add)	(None, 2, 2, 96)	0	block_10_project... block_11_project...
block_12_expand (Conv2D)	(None, 2, 2, 576)	55,296	block_11_add[0][...
block_12_expand_BN (BatchNormalizatio...	(None, 2, 2, 576)	2,304	block_12_expand[...
block_12_expand_re... (ReLU)	(None, 2, 2, 576)	0	block_12_expand_...

block_12_depthwise (DepthwiseConv2D)	(None, 2, 2, 576)	5,184	block_12_expand_...
block_12_depthwise... (BatchNormalizatio...	(None, 2, 2, 576)	2,304	block_12_depthwi...
block_12_depthwise... (ReLU)	(None, 2, 2, 576)	0	block_12_depthwi...
block_12_project (Conv2D)	(None, 2, 2, 96)	55,296	block_12_depthwi...
block_12_project_BN (BatchNormalizatio...	(None, 2, 2, 96)	384	block_12_project...
block_12_add (Add)	(None, 2, 2, 96)	0	block_11_add[0] [...] block_12_project...
block_13_expand (Conv2D)	(None, 2, 2, 576)	55,296	block_12_add[0] [...]
block_13_expand_BN (BatchNormalizatio...	(None, 2, 2, 576)	2,304	block_13_expand[...
block_13_expand_re... (ReLU)	(None, 2, 2, 576)	0	block_13_expand_...
block_13_pad (ZeroPadding2D)	(None, 3, 3, 576)	0	block_13_expand_...
block_13_depthwise (DepthwiseConv2D)	(None, 1, 1, 576)	5,184	block_13_pad[0] [...]
block_13_depthwise... (BatchNormalizatio...	(None, 1, 1, 576)	2,304	block_13_depthwi...
block_13_depthwise... (ReLU)	(None, 1, 1, 576)	0	block_13_depthwi...
block_13_project (Conv2D)	(None, 1, 1, 160)	92,160	block_13_depthwi...
block_13_project_BN (BatchNormalizatio...	(None, 1, 1, 160)	640	block_13_project...
block_14_expand (Conv2D)	(None, 1, 1, 960)	153,600	block_13_project...

block_14_expand_BN (BatchNormalizatio...	(None, 1, 1, 960)	3,840	block_14_expand[...
block_14_expand_re... (ReLU)	(None, 1, 1, 960)	0	block_14_expand_...
block_14_depthwise (DepthwiseConv2D)	(None, 1, 1, 960)	8,640	block_14_expand_...
block_14_depthwise... (BatchNormalizatio...	(None, 1, 1, 960)	3,840	block_14_depthwi...
block_14_depthwise... (ReLU)	(None, 1, 1, 960)	0	block_14_depthwi...
block_14_project (Conv2D)	(None, 1, 1, 160)	153,600	block_14_depthwi...
block_14_project_BN (BatchNormalizatio...	(None, 1, 1, 160)	640	block_14_project...
block_14_add (Add)	(None, 1, 1, 160)	0	block_13_project... block_14_project...
block_15_expand (Conv2D)	(None, 1, 1, 960)	153,600	block_14_add[0][...
block_15_expand_BN (BatchNormalizatio...	(None, 1, 1, 960)	3,840	block_15_expand[...
block_15_expand_re... (ReLU)	(None, 1, 1, 960)	0	block_15_expand_...
block_15_depthwise (DepthwiseConv2D)	(None, 1, 1, 960)	8,640	block_15_expand_...
block_15_depthwise... (BatchNormalizatio...	(None, 1, 1, 960)	3,840	block_15_depthwi...
block_15_depthwise... (ReLU)	(None, 1, 1, 960)	0	block_15_depthwi...
block_15_project (Conv2D)	(None, 1, 1, 160)	153,600	block_15_depthwi...
block_15_project_BN (BatchNormalizatio...	(None, 1, 1, 160)	640	block_15_project...

block_15_add (Add)	(None, 1, 1, 160)	0	block_14_add[0][...] block_15_project...
block_16_expand (Conv2D)	(None, 1, 1, 960)	153,600	block_15_add[0][...]
block_16_expand_BN (BatchNormalizatio...	(None, 1, 1, 960)	3,840	block_16_expand[...]
block_16_expand_re... (ReLU)	(None, 1, 1, 960)	0	block_16_expand_...
block_16_depthwise (DepthwiseConv2D)	(None, 1, 1, 960)	8,640	block_16_expand_...
block_16_depthwise... (BatchNormalizatio...	(None, 1, 1, 960)	3,840	block_16_depthwi...
block_16_depthwise... (ReLU)	(None, 1, 1, 960)	0	block_16_depthwi...
block_16_project (Conv2D)	(None, 1, 1, 320)	307,200	block_16_depthwi...
block_16_project_BN (BatchNormalizatio...	(None, 1, 1, 320)	1,280	block_16_project...
Conv_1 (Conv2D)	(None, 1, 1, 1280)	409,600	block_16_project...
Conv_1_bn (BatchNormalizatio...	(None, 1, 1, 1280)	5,120	Conv_1[0][0]
out_relu (ReLU)	(None, 1, 1, 1280)	0	Conv_1_bn[0][0]
global_average_poo... (GlobalAveragePool...	(None, 1280)	0	out_relu[0][0]
dense_2 (Dense)	(None, 128)	163,968	global_average_p...
dropout_1 (Dropout)	(None, 128)	0	dense_2[0][0]
dense_3 (Dense)	(None, 43)	5,547	dropout_1[0][0]

Total params: 2,427,499 (9.26 MB)

Trainable params: 169,515 (662.17 KB)

Non-trainable params: 2,257,984 (8.61 MB)

Train MobileNetV2

```
[ ]: history_mobilenet = mobilenet_model.fit(
    datagen.flow(X_train, y_train, batch_size=32),
    validation_data=(X_val, y_val),
    epochs=7
)
```

Epoch 1/7

981/981 27s 28ms/step -

accuracy: 0.2825 - loss: 2.4340 - val\_accuracy: 0.3192 - val\_loss: 2.2776

Epoch 2/7

981/981 28s 28ms/step -

accuracy: 0.2860 - loss: 2.3973 - val\_accuracy: 0.3277 - val\_loss: 2.2460

Epoch 3/7

981/981 27s 28ms/step -

accuracy: 0.2932 - loss: 2.3779 - val\_accuracy: 0.3315 - val\_loss: 2.2198

Epoch 4/7

981/981 26s 26ms/step -

accuracy: 0.2984 - loss: 2.3605 - val\_accuracy: 0.3333 - val\_loss: 2.2042

Epoch 5/7

981/981 41s 26ms/step -

accuracy: 0.2972 - loss: 2.3432 - val\_accuracy: 0.3391 - val\_loss: 2.1873

Epoch 6/7

981/981 24s 25ms/step -

accuracy: 0.3020 - loss: 2.3190 - val\_accuracy: 0.3406 - val\_loss: 2.1776

Epoch 7/7

981/981 24s 24ms/step -

accuracy: 0.3061 - loss: 2.3096 - val\_accuracy: 0.3486 - val\_loss: 2.1603


Compare Models

```
[ ]: import matplotlib.pyplot as plt

def plot_history(histories, titles):
    plt.figure(figsize=(12,5))
    for history, label in zip(histories, titles):
        plt.plot(history.history['val_accuracy'], label=f'{label} Val Acc')
    plt.title('Validation Accuracy Comparison')
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy')
```

```
plt.legend()
plt.show()

plot_history(
    [history, history_aug, history_mobilenet],
    ['Custom CNN', 'CNN + Augmentation', 'MobileNetV2']
)
```



```
[ ]: %%shell
cd Elevvo-Pathways-ML-intern/
git add .
git commit -m "Update notebook with tasks"
git push origin HEAD
```